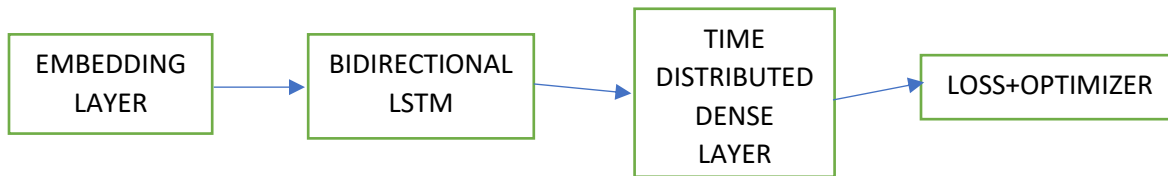


Model used (Attached a link at the end of document for model given by tensorboard):



```

embeddings = tf.get_variable('embeddings', [self.num_terms, 30])
model_output = tf.nn.embedding_lookup(embeddings, self.x)
cell = tf.keras.layers.LSTMCell(48)
layer2 = tf.keras.layers.Bidirectional(tf.keras.layers.RNN(cell, return_sequences=True))
model_output = layer2(model_output)
layer3 = tf.keras.layers.TimeDistributed(tf.keras.layers.Dense(self.num_tags))
self.logits = layer3(model_output)
  
```

Stuff that improved my model:

Hyper parameter tuning:

Learning rate (LR) is one of the most important hyper parameter that affects any model as the time for each language data set is quite small. Intuitively, I started with a high LR and decreased LR after each iteration as follows:

Initial learn_rate = 4×10^{-2}

learn_rate *= ((0.8)**(self.iter)) #new learn_rate = (initial learn rate)*(0.8)^{iteration number - 1}

learn_rate /= (self.iter+1) #new learn rate = new_learn_rate/(iteration number)

This is done because we want to decrease the loss fast early and slower in later iterations.

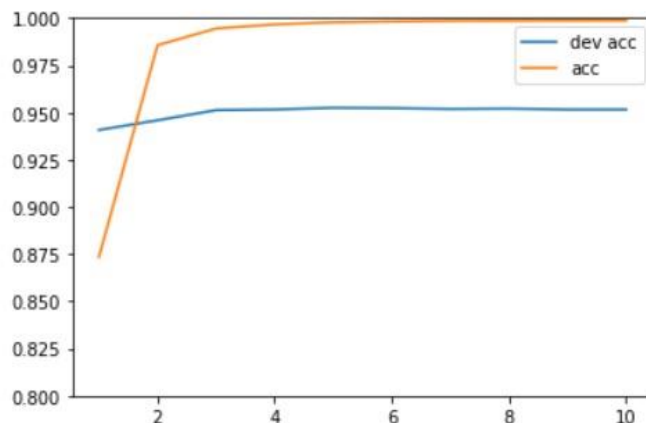
Embedding dimension size: I tried with many sizes but 30 worked best for both languages

Number of units for LSTM Cell: 48 units gave the best output for both languages

Early stopping: I used max iteration as 4 -> So it will run at most 4 iterations for any language. I observed after this, we are over fitting the model/dev accuracy doesn't change much after 4 iterations (as shown below).

Batch size: I used a batch size of 64.

Dropout: I tried different dropout and recurrent dropout, but it didn't have much effect.



Training vs Testing Accuracies for Japanese language

Loss used:

I experimented with different losses.

- 1) seq2seq.sequence_loss + Adam Optimizer

```
losses = tf.contrib.seq2seq.sequence_loss(self.logits,self.target,self.weights)
opt = tf.train.AdamOptimizer(learning_rate=self.learning_rate)
self.train_op = opt.minimize(losses)

#predict tags as or decoding as:
numpy.argmax(logits, axis=2)
```

- 2) CRF - https://en.wikipedia.org/wiki/Conditional_random_field + Adam Optimizer

```
crf_params = tf.get_variable("CRF", [self.num_tags, self.num_tags], dtype=tf.float32)
log_likelihood, self.transition_params = tf.contrib.crf.crf_log_likelihood(
    self.logits,self.target,self.lengths,crf_params)
losses = tf.reduce_mean(-log_likelihood)
opt = tf.train.AdamOptimizer(learning_rate=self.learning_rate)
self.train_op = opt.minimize(losses)
```

and predict tags as (Viterbi Decoding):

```
logits,transitions = self.sess.run([self.logits,self.transition_params], {self.x: terms, self.lengths: lengths})
#get logits and transition matrix
pred_tags = numpy.zeros(terms.shape) #create a matrix for predicted tags
for seq_len, logits_data in zip(lengths,logits):
    score = numpy.array(logits_data[:seq_len]).astype(numpy.float32)
    viterbi_sequence, viterbi_score = tf.contrib.crf.viterbi_decode(
        score, transitions) #for each sentence, get the corresponding score from logits and use
    the transition for Viterbi decoding to get the sequence of tags
    pred_tags[i,:seq_len] = viterbi_sequence
    i+=1
return pred_tags
```

I tried with both 1 and 2, they gave almost the same results but went with CRF because thought it will fit better with surprise language.

Doing the above, I got an accuracy of **~95.7%** on **Italian**, **~95.2%** on **Japanese** and **>95th** Percentile on the **Surprise Language**.

Model as given by tensorboard:

CRF Model: https://drive.google.com/file/d/1SYaqOQoJW_Nf_9pJ46YPJreLOQENF9oQ/view -> looks complicated but TensorFlow does most of the work.

Normal model: https://drive.google.com/open?id=1KF1IjXdvzoW_qgSag-BVw8vtronouD11F